

# Server-side Data Reduction and Analysis with Script Workflow Analysis for MultiProcessing (SWAMP)

Daniel L. Wang, Charles S. Zender, and Stephen F. Jenks  
University of California, Irvine

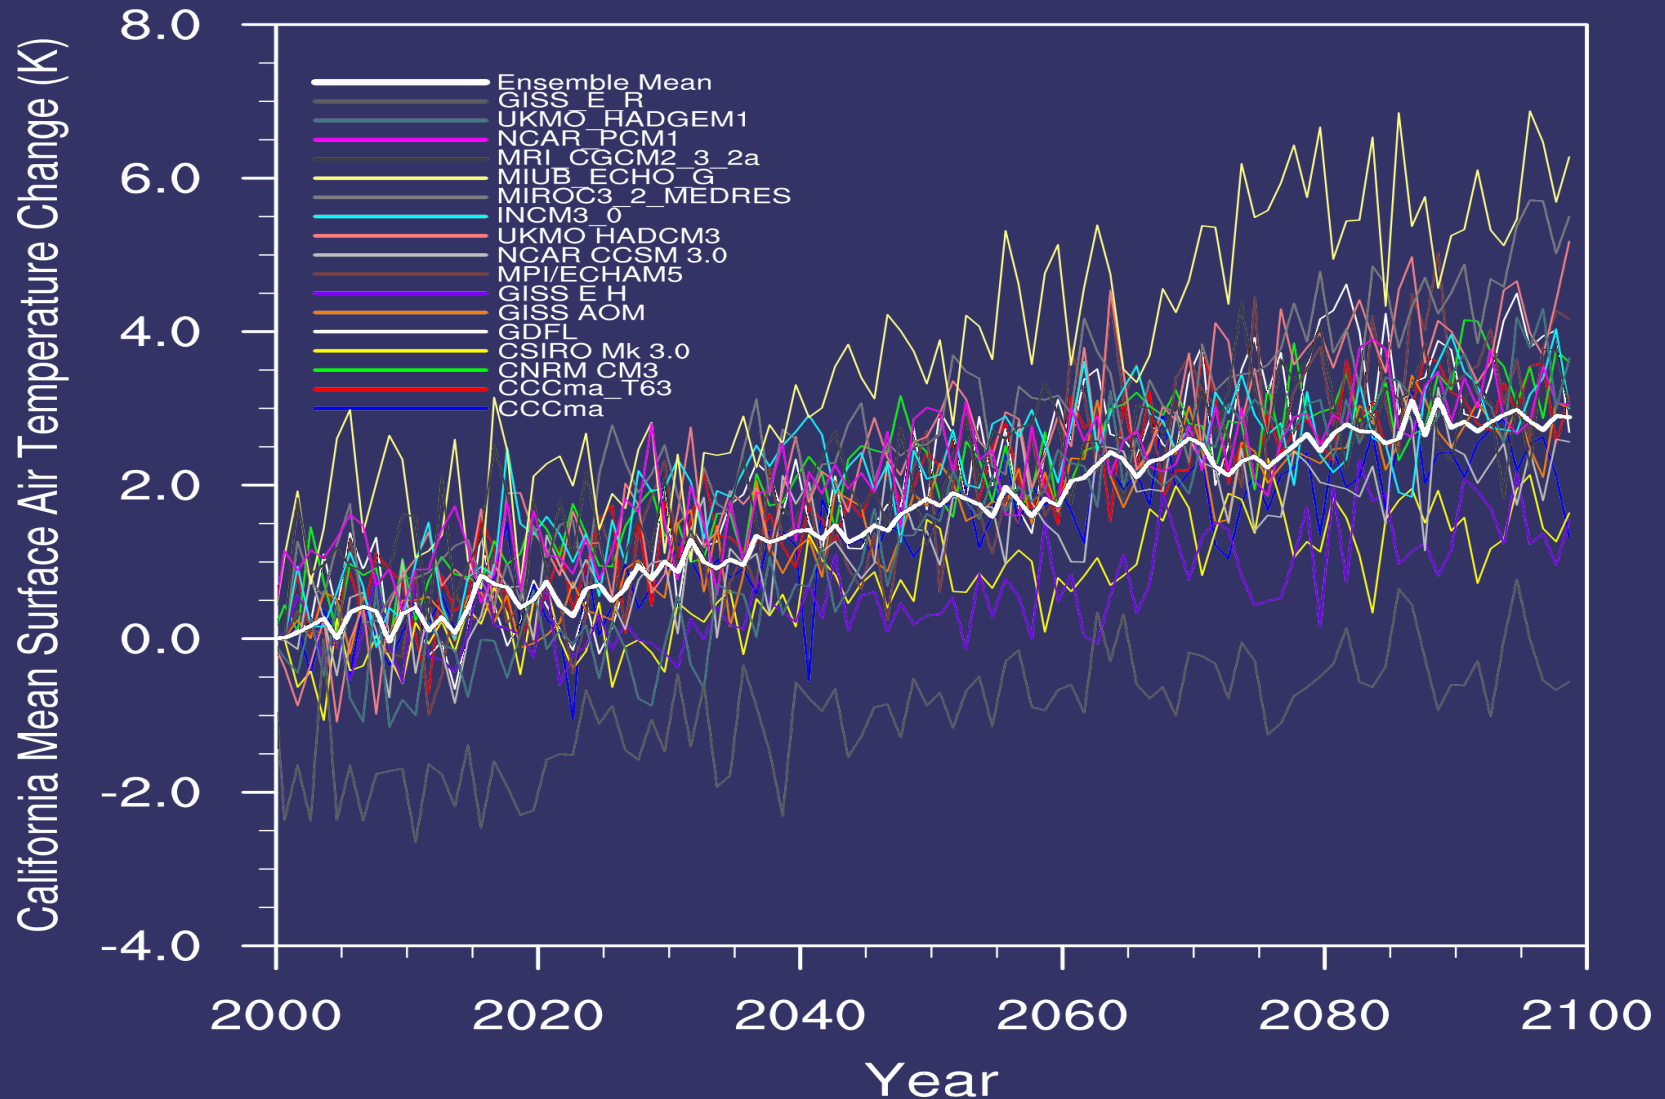
[http://dust.ess.uci.edu/~wangd/pub/sld\\_WZJ07b.pdf](http://dust.ess.uci.edu/~wangd/pub/sld_WZJ07b.pdf)



# Motivation

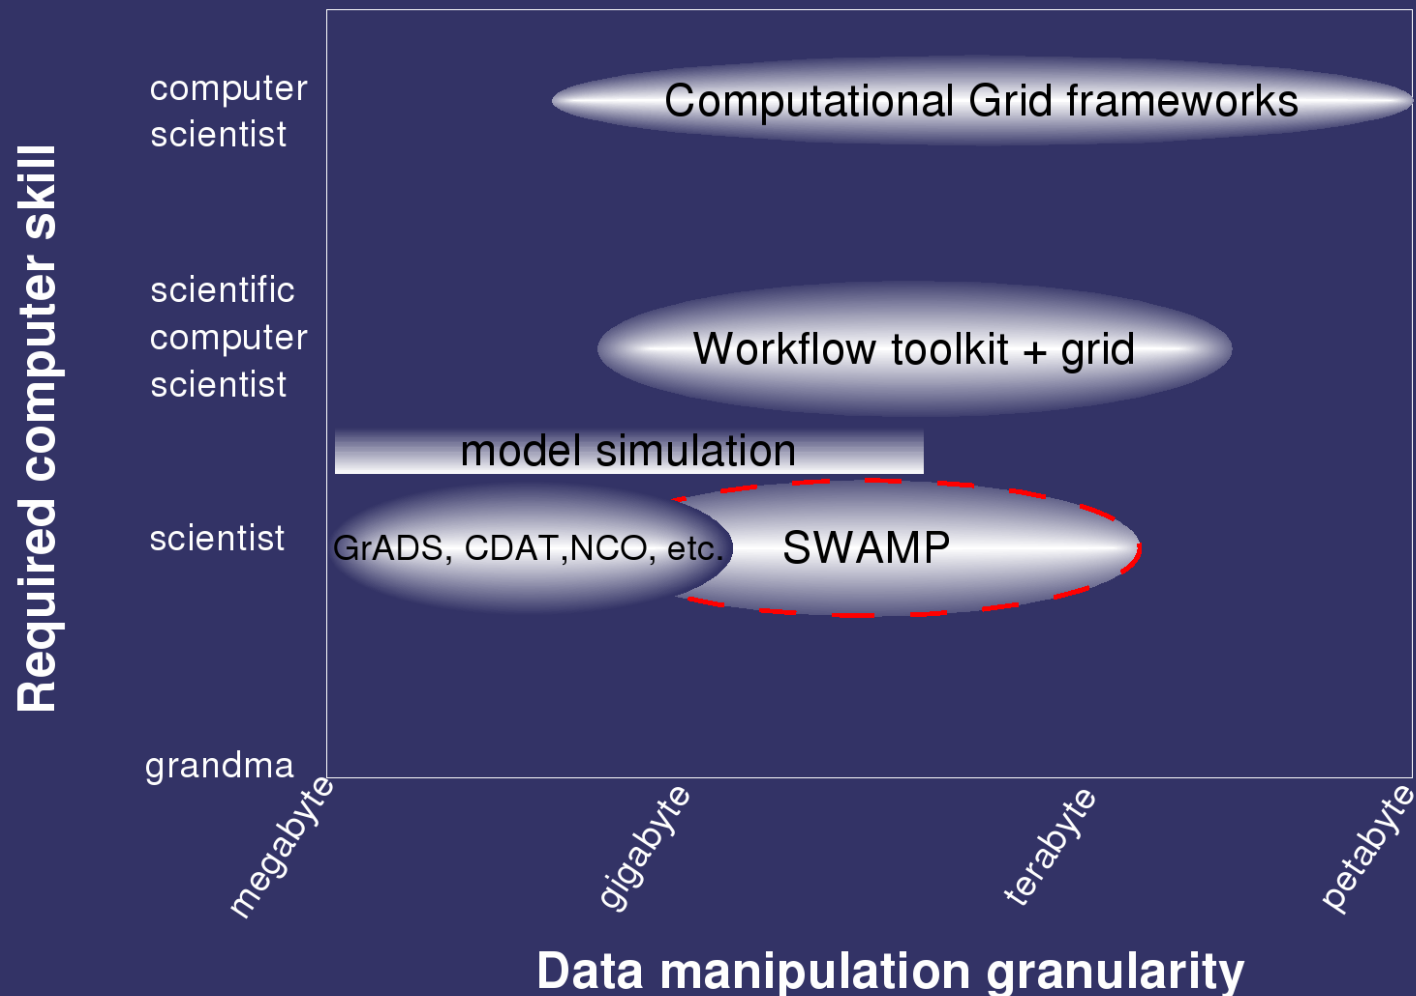
➡ How can we make this task easy?

## SRESA1B 720ppm Stabilization Scenario

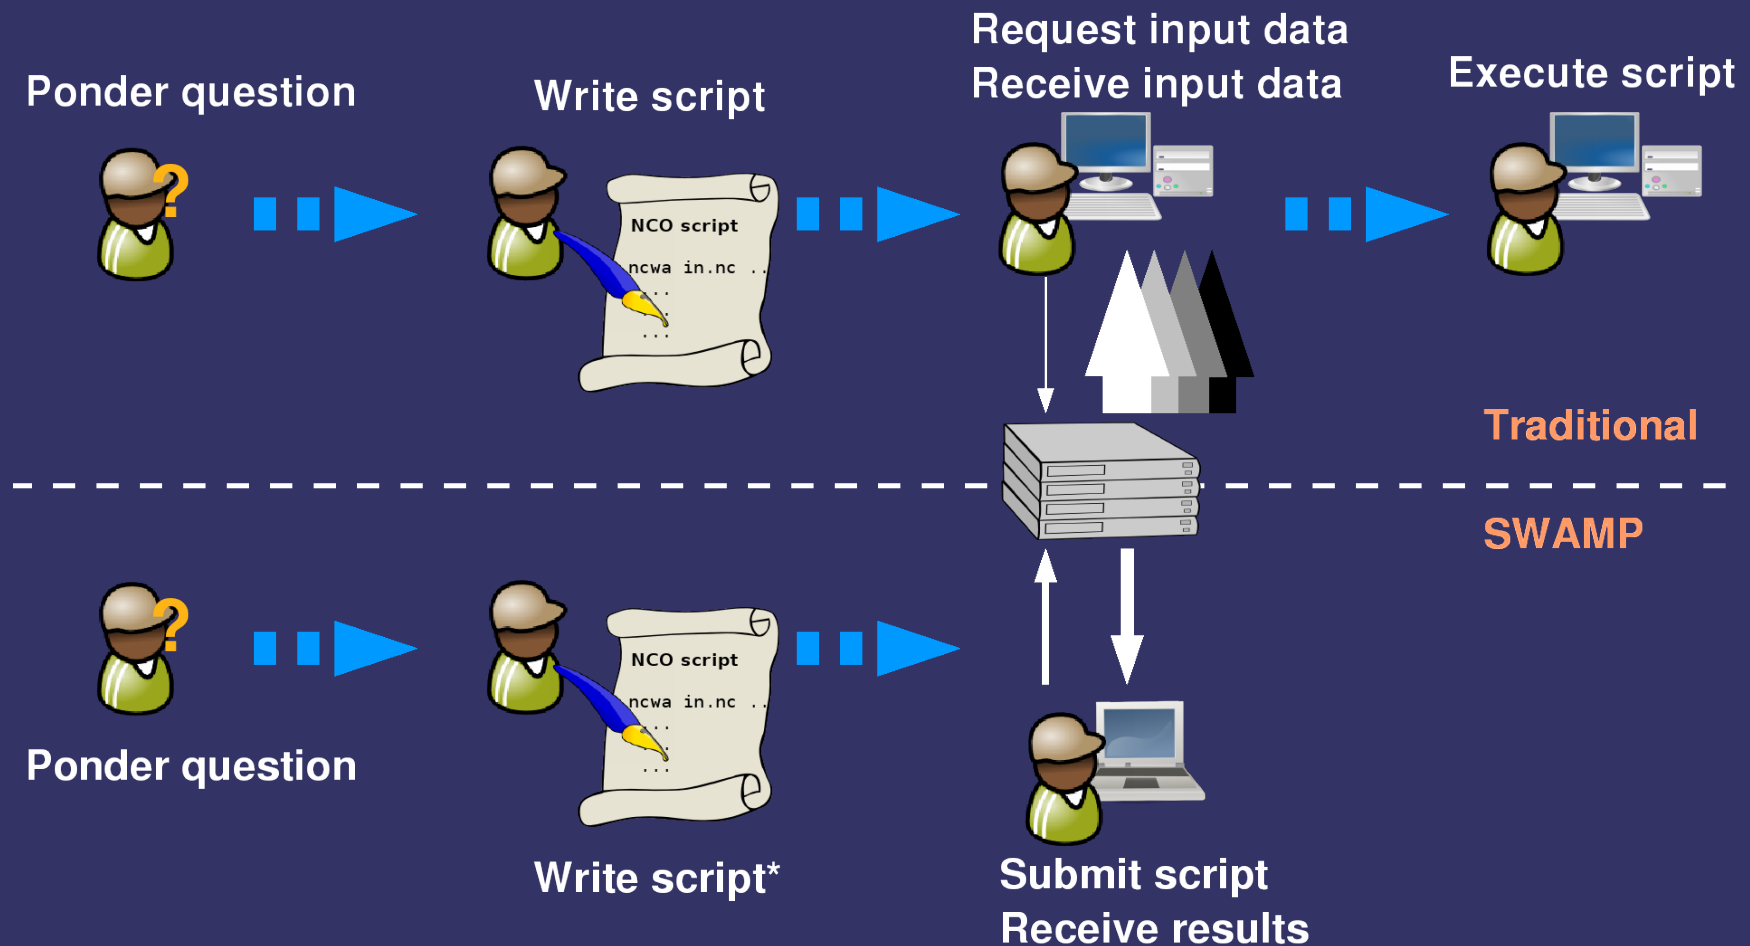


# Problem

- ➔ Analysis lags simulation/generation
  - Bulk: 1TB of data is non-trivial to manage
  - Locality: Bandwidth  $\ll \infty$



# Approach



- ⇒ Appropriately distribute computation load
- ⇒ Leverage existing scripted interfaces
- ⇒ (Exploit available parallelism)

# *The “Normal” Way*

```
models='cccma_cgcm3_1 cccma_cgcm3_1_t63 cnrm_cm3 csiro_mk3_0 \  
      gfdl_cm2_0 gfdl_cm2_1 giss_aom giss_model_e_h giss_model_e_r \  
      iap_fgoals1_0_g inmcm3_0 ipsl_cm4 miroc3_2_hires miroc3_2_medres \  
      miub_echo_g mpi_echam5 mri_cgcm2_3_2a ncar_ccsm3_0 ncar_pcm1 \  
      ukmo_hadcm3 ukmo_hadgem1'  
variables='tas pr'  
scenarios='sresalb sresa2 sresb1'  
  
for scn in $scenarios; do  
  
    for mdl in $models; do  
  
        ncwa -O -v $variables -w area -a lat,lon \  
            -p http://user:password@climate.llnl.gov/cgi-bin/dap-cgi.py/ipcc4/$scn/$mdl \  
            pcmdi.ipcc4.$mdl.$scn.run1.atm.mo.xml $scn_$mdl_200001_209912.nc  
        ncwa -F -d time,1,12 $scn_$mdl_200001_209912.nc $scn_$mdl_2000.nc  
        ncdiff $scn_$mdl_200001_209912.nc $scn_$mdl_2000.nc $scn_$mdl_anm.nc  
  
    done # end loop over model  
  
    ncea *_200001_209912.nc $scn_avg_200001_209912.nc  
    ncwa -F -d time,1,12 $scn_avg_200001_209912.nc $scn_avg_2000.nc  
    ncdiff $scn_avg_200001_209912.nc $scn_avg_2000.nc $scn_avg_anm.nc  
  
done # end loop over scenario
```

# SWAMP

## Script Workflow Analysis for Multiprocessing

- ➔ Specialized netCDF data handler for OPeNDAP
  - Use constraint expression to signal computation
- ➔ ~sh-script syntax accepts netCDF Operators (NCO) command lines
  - `ncwa -v WIND -d time %tempf_cam1999.nc%\  
%outf_winds1999.nc%`
- ➔ Server-side parallelizing and optimizing execution engine

# *Engine implementation*

- ⇒ Replaces the netCDF data handler (dap\_nc\_handler)
  - Passes non-SWAMP requests to original
- ⇒ Written in python
- ⇒ Parses script to generate workflow
- ⇒ Locates temporary files in ramdisk, 'live file' analysis to reclaim space
- ⇒ “Persistent” job tracking via database
  - SQLite for speed and simplicity

# *Engine Parallelization*

- ⇒ Peer worker parallelization model
  - server-configurable # of slots
  - shared state via DB updates
  - ready list, file status, command state
- ⇒ Script-line-level parallelization
  - thread-level built-in to NCO (OpenMP)
- ⇒ I/O concurrency a problem for  $n > 2$ 
  - Use Linux tmpfs ramdisk



# *Experimental Setup*

⇒ Case: Subsample 10 years of global T42 data at  $\Delta t=20\text{m}$  into  $\Delta t=12\text{h}$

- ~ 14,000 script lines
- ~ 8GB input data (120 files)
- ~ 26GB intermediate results
- ~ 230MB result data (10 files)

230MB



8GB



26GB

⇒ System:

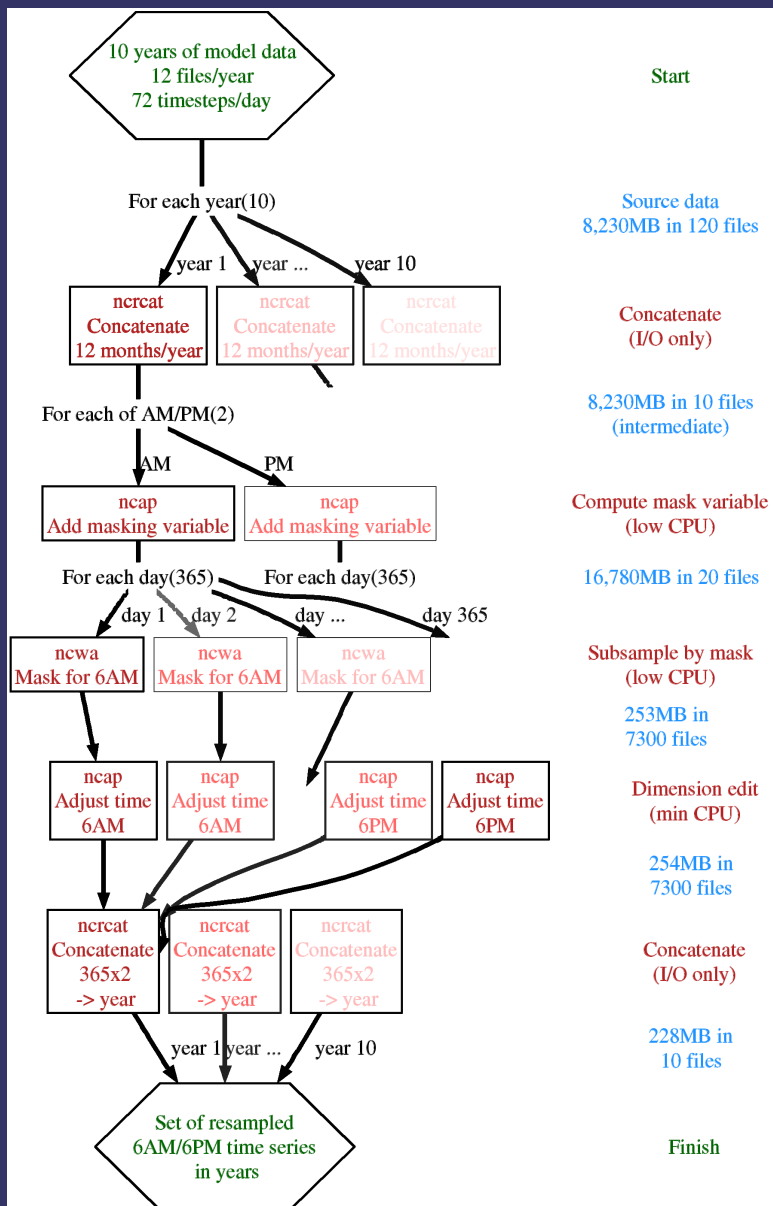
- Dual Opteron 270 (4 cores total), 16GB memory

⇒ Compare 9 cases:

- Traditional vs. SWAMP (1,2,4,8-wide, I/O opt)




# SWAMP Parses, Optimizes and Schedules

[illegible]

~14,000 line script  
1 line ~ 1 NCO command

## dependency-aware workflow

# SWAMP Performance



# SWAMP Parallelization



# *Conclusions*

- ⇒ High computational performance potential with tested workflow
  - Parallelism
  - Optimization
- ⇒ Large speedup through bandwidth optimization
- ⇒ Viability of shell-script interface
- ⇒ Room for optimization above existing NCO optimization

# *Would you like to know more?*

## Questions?

- offline: wangd@uci.edu

## Want to poke at the code?

(please be gentle)

- See NCO CVS source tree on sourceforge:  
( nco/src/ssdap )

## Want to try it?

- Probably need my help to make it work